

```
In [1]: # Assignment: DSC550 Week10
        # Name: Bezawada, Sashidhar
        # Date: 2023-05-20
        # 8.2 Term Project: Term Project Milestone 3: Model Building and Evaluation
```

```
In [2]: #import Libraries
        import numpy as np
        import pandas as pd
        import matplotlib.pyplot as plt
```

```
In [3]: #Load the dataset as a Pandas data frame
        df_brainstroke= pd.read_csv("datasets/week6/brain_stroke.csv")
```

```
In [4]: #Display the first ten rows of data
        df_brainstroke.head(10)
```

```
Out[4]:
```

	id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_g
0	9046	Male	67.0	0	1	Yes	Private	Urban	
1	51676	Female	61.0	0	0	Yes	Self-employed	Rural	
2	31112	Male	80.0	0	1	Yes	Private	Rural	
3	60182	Female	49.0	0	0	Yes	Private	Urban	
4	1665	Female	79.0	1	0	Yes	Self-employed	Rural	
5	56669	Male	81.0	0	0	Yes	Private	Urban	
6	53882	Male	74.0	1	1	Yes	Private	Rural	
7	10434	Female	69.0	0	0	No	Private	Urban	
8	27419	Female	59.0	0	0	Yes	Private	Rural	
9	60491	Female	78.0	0	0	Yes	Private	Urban	

```
In [5]: #Find the information in the data frame.
        df_brainstroke.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5110 entries, 0 to 5109
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     5110 non-null   int64
1   gender                 5110 non-null   object
2   age                   5110 non-null   float64
3   hypertension           5110 non-null   int64
4   heart_disease          5110 non-null   int64
5   ever_married           5110 non-null   object
6   work_type              5110 non-null   object
7   Residence_type         5110 non-null   object
8   avg_glucose_level      5110 non-null   float64
9   bmi                    4909 non-null   float64
10  smoking_status         5110 non-null   object
11  stroke                 5110 non-null   int64
dtypes: float64(3), int64(4), object(5)
memory usage: 479.2+ KB
In [6]: df_brainstroke['gender'].unique()

Out[6]:array(['Male', 'Female', 'Other'], dtype=object)
In [7]: df_brainstroke['hypertension'].unique()

Out[7]:array([0, 1], dtype=int64)
In [8]: df_brainstroke['heart_disease'].unique()

Out[8]:array([1, 0], dtype=int64)
In [9]: df_brainstroke['ever_married'].unique()

Out[9]:array(['Yes', 'No'], dtype=object)
In [10]: df_brainstroke['work_type'].unique()

Out[10]:array(['Private', 'Self-employed', 'Govt_job', 'children', 'Never_worked'],
              dtype=object)
In [11]: df_brainstroke['stroke'].unique()

Out[11]:array([1, 0], dtype=int64)

```

MILESTONE 2

Data Transformation

1. Dropping Columns/Features

Dropping column Id as this a unique number given to each record (can be patient id or random generated number or serial number). It has no significance in decision making.

```
In [12]: df_brainstroke = df_brainstroke.drop('id',axis=1)
```

2. Replacing values in a cloumn

Changing work type 'Children' as 'Student' as Children is not work type and Changing formatting for other work types. Changing binary values to string values of Yes and No instead of 1 and 0 respectively for hypertension and heart_disease.

```

In [14]: df_brainstroke['work_type'].unique()

Out[14]: array(['Private', 'Self Employed', 'Government', 'Student',
               'Never_worked'], dtype=object)

In ... #Changing binary values to string values of Yes and No instead of 1 and 0 respect
df_brainstroke['hypertension'] = df_brainstroke['hypertension'].replace([1], 'Yes')
df_brainstroke['hypertension'] = df_brainstroke['hypertension'].replace([0], 'No')

df_brainstroke['heart_disease'] = df_brainstroke['heart_disease'].replace([1], 'Y')
df_brainstroke['heart_disease'] = df_brainstroke['heart_disease'].replace([0], 'N')

In [16]: df_brainstroke['hypertension'].unique()

Out[16]: array(['No', 'Yes'], dtype=object)

In [17]: df_brainstroke['heart_disease'].unique()

Out[17]: array(['Yes', 'No'], dtype=object)

In [18]: df_brainstroke['stroke'].unique()

Out[18]: array([1, 0], dtype=int64)

In [19]: df_brainstroke['smoking_status'].unique()

Out[19]: array(['formerly smoked', 'never smoked', 'smokes', 'Unknown'],
               dtype=object)

```

3. Renaming Column Name

Changing column name stroke to brain_stroke as this is more meaningful

```

In [20]: df_brainstroke = df_brainstroke.rename(columns={'stroke': 'brain_stroke'})

In [21]: df_brainstroke.head(10)

```

```

Out[21]:

```

	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose
0	Male	67.0	No	Yes	Yes	Private	Urban	
1	Female	61.0	No	No	Yes	Self Employed	Rural	
2	Male	80.0	No	Yes	Yes	Private	Rural	
3	Female	49.0	No	No	Yes	Private	Urban	
4	Female	79.0	Yes	No	Yes	Self Employed	Rural	
5	Male	81.0	No	No	Yes	Private	Urban	
6	Male	74.0	Yes	Yes	Yes	Private	Rural	
7	Female	69.0	No	No	No	Private	Urban	
8	Female	59.0	No	No	Yes	Private	Rural	
9	Female	78.0	No	No	Yes	Private	Urban	

```
In [22]: df_brainstroke.dtypes
```

```
Out[22]:gender                object
age                float64
hypertension       object
heart_disease      object
ever_married       object
work_type          object
Residence_type     object
avg_glucose_level  float64
bmi                float64
smoking_status     object
brain_stroke       int64
dtype: object
```

4. Transform features

As per BMI chart on National Heart, Lung and Blood Institute website (https://www.nhlbi.nih.gov/health/educational/lose_wt/BMI/bmi_tbl2.html), the maximum BMI is 54. So, any value above 54 is unrealistic. So, replace values over 54 with 54 considering the person is obese.

```
In [23]: np.max(df_brainstroke['bmi'])
```

```
Out[23]:97.6
```

```
In [24]: max_bmi = 54.0
         df_brainstroke['bmi'][df_brainstroke['bmi']>max_bmi] = max_bmi
```

C:\Users\sashi_000\AppData\Local\Temp\ipykernel_15068\849835967.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_brainstroke['bmi'][df_brainstroke['bmi']>max_bmi] = max_bmi
```

```
In [25]: np.max(df_brainstroke['bmi'])
```

```
Out[25]:54.0
```

5. Engineer new useful features.

Adding new column stage_of_life based on the age as we can decide which age group has an impact

```
In [... df_brainstroke['stage_of_life'] = np.where(df_brainstroke.age <= 2, 'Infant',
np.where(df_brainstroke.age < 5, 'Toddler',
np.where(df_brainstroke.age < 13, 'Child',
np.where(df_brainstroke.age < 20, 'Teen',
np.where(df_brainstroke.age < 40, 'Adult',
np.where(df_brainstroke.age < 60, 'Middle Age A
```

```
In [27]: df_brainstroke.head(10)
```

Out[27]:

	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose
0	Male	67.0	No	Yes	Yes	Private	Urban	2
1	Female	61.0	No	No	Yes	Self Employed	Rural	2
2	Male	80.0	No	Yes	Yes	Private	Rural	1
3	Female	49.0	No	No	Yes	Private	Urban	1
4	Female	79.0	Yes	No	Yes	Self Employed	Rural	1
5	Male	81.0	No	No	Yes	Private	Urban	1
6	Male	74.0	Yes	Yes	Yes	Private	Rural	
7	Female	69.0	No	No	No	Private	Urban	
8	Female	59.0	No	No	Yes	Private	Rural	
9	Female	78.0	No	No	Yes	Private	Urban	

< >

I am adding one more column 'weight_status' as bmi might have different values and difficult guage details on that. So, weight_status will categorize them into 3 different categories. Instead of bmi, I feel weight_status will be best suited for analysis.

```
In [... df_brainstroke['weight_status'] = np.where(df_brainstroke.bmi < 18.5, 'Underweig
np.where(df_brainstroke.bmi < 25.0, 'Healthy we
np.where(df_brainstroke.bmi < 30, 'Overweight',
```

```
In [29]: df_brainstroke.head(10)
```

```
Out[29]:
```

	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose
0	Male	67.0	No	Yes	Yes	Private	Urban	2
1	Female	61.0	No	No	Yes	Self Employed	Rural	2
2	Male	80.0	No	Yes	Yes	Private	Rural	1
3	Female	49.0	No	No	Yes	Private	Urban	1
4	Female	79.0	Yes	No	Yes	Self Employed	Rural	1
5	Male	81.0	No	No	Yes	Private	Urban	1
6	Male	74.0	Yes	Yes	Yes	Private	Rural	
7	Female	69.0	No	No	No	Private	Urban	
8	Female	59.0	No	No	Yes	Private	Rural	
9	Female	78.0	No	No	Yes	Private	Urban	

6. Replace Null Values

```
In [30]: df_brainstroke.isnull().sum()
```

```
Out[30]:gender          0
age                    0
hypertension          0
heart_disease         0
ever_married          0
work_type             0
Residence_type        0
avg_glucose_level     0
bmi                   201
smoking_status        0
brain_stroke          0
stage_of_life         0
weight_status         0
dtype: int64
```

As there are 201 null values for bmi, if we remove them, it will reduce dataset drastically. So, I am replacing null values with median.

```
In [31]: df_brainstroke['bmi'] = df_brainstroke['bmi'].fillna(df_brainstroke['bmi'].median())
```

```
In [32]: df_brainstroke.isnull().sum()
```

```
Out[32]:gender      0
        age         0
        hypertension 0
        heart_disease 0
        ever_married 0
        work_type    0
        Residence_type 0
        avg_glucose_level 0
        bmi          0
        smoking_status 0
        brain_stroke 0
        stage_of_life 0
        weight_status 0
        dtype: int64
```

7. String Values case change

Finally, changing string values to UPPER CASE to avoid same values due to different cases falling into 2 different categories

```
In [33]: df_brainstroke = df_brainstroke.apply(lambda x: x.astype(str).str.upper())
```

```
In [34]: np.sort(df_brainstroke['bmi'].unique())
```

```

Out[34]:array([ '10.3', '11.3', '11.5', '12.0', '12.3', '12.8', '13.0', '13.2',
'13.3', '13.4', '13.5', '13.7', '13.8', '13.9', '14.0', '14.1',
'14.2', '14.3', '14.4', '14.5', '14.6', '14.8', '14.9', '15.0',
'15.1', '15.2', '15.3', '15.4', '15.5', '15.6', '15.7', '15.8',
'15.9', '16.0', '16.1', '16.2', '16.3', '16.4', '16.5', '16.6',
'16.7', '16.8', '16.9', '17.0', '17.1', '17.2', '17.3', '17.4',
'17.5', '17.6', '17.7', '17.8', '17.9', '18.0', '18.1', '18.2',
'18.3', '18.4', '18.5', '18.6', '18.7', '18.8', '18.9', '19.0',
'19.1', '19.2', '19.3', '19.4', '19.5', '19.6', '19.7', '19.8',
'19.9', '20.0', '20.1', '20.2', '20.3', '20.4', '20.5', '20.6',
'20.7', '20.8', '20.9', '21.0', '21.1', '21.2', '21.3', '21.4',
'21.5', '21.6', '21.7', '21.8', '21.9', '22.0', '22.1', '22.2',
'22.3', '22.4', '22.5', '22.6', '22.7', '22.8', '22.9', '23.0',
'23.1', '23.2', '23.3', '23.4', '23.5', '23.6', '23.7', '23.8',
'23.9', '24.0', '24.1', '24.2', '24.3', '24.4', '24.5', '24.6',
'24.7', '24.8', '24.9', '25.0', '25.1', '25.2', '25.3', '25.4',
'25.5', '25.6', '25.7', '25.8', '25.9', '26.0', '26.1', '26.2',
'26.3', '26.4', '26.5', '26.6', '26.7', '26.8', '26.9', '27.0',
'27.1', '27.2', '27.3', '27.4', '27.5', '27.6', '27.7', '27.8',
'27.9', '28.0', '28.1', '28.2', '28.3', '28.4', '28.5', '28.6',
'28.7', '28.8', '28.9', '29.0', '29.1', '29.2', '29.3', '29.4',
'29.5', '29.6', '29.7', '29.8', '29.9', '30.0', '30.1', '30.2',
'30.3', '30.4', '30.5', '30.6', '30.7', '30.8', '30.9', '31.0',
'31.1', '31.2', '31.3', '31.4', '31.5', '31.6', '31.7', '31.8',
'31.9', '32.0', '32.1', '32.2', '32.3', '32.4', '32.5', '32.6',
'32.7', '32.8', '32.9', '33.0', '33.1', '33.2', '33.3', '33.4',
'33.5', '33.6', '33.7', '33.8', '33.9', '34.0', '34.1', '34.2',
'34.3', '34.4', '34.5', '34.6', '34.7', '34.8', '34.9', '35.0',
'35.1', '35.2', '35.3', '35.4', '35.5', '35.6', '35.7', '35.8',
'35.9', '36.0', '36.1', '36.2', '36.3', '36.4', '36.5', '36.6',
'36.7', '36.8', '36.9', '37.0', '37.1', '37.2', '37.3', '37.4',
'37.5', '37.6', '37.7', '37.8', '37.9', '38.0', '38.1', '38.2',
'38.3', '38.4', '38.5', '38.6', '38.7', '38.8', '38.9', '39.0',
'39.1', '39.2', '39.3', '39.4', '39.5', '39.6', '39.7', '39.8',
'39.9', '40.0', '40.1', '40.2', '40.3', '40.4', '40.5', '40.6',
'40.7', '40.8', '40.9', '41.0', '41.1', '41.2', '41.3', '41.4',
'41.5', '41.6', '41.7', '41.8', '41.9', '42.0', '42.1', '42.2',
'42.3', '42.4', '42.5', '42.6', '42.7', '42.8', '42.9', '43.0',
'43.1', '43.2', '43.3', '43.4', '43.6', '43.7', '43.8', '43.9',
'44.0', '44.1', '44.2', '44.3', '44.4', '44.5', '44.6', '44.7',
'44.8', '44.9', '45.0', '45.1', '45.2', '45.3', '45.4', '45.5',
'45.7', '45.8', '45.9', '46.0', '46.1', '46.2', '46.3', '46.4',
'46.5', '46.6', '46.8', '46.9', '47.1', '47.3', '47.4', '47.5',
'47.6', '47.8', '47.9', '48.0', '48.1', '48.2', '48.3', '48.4',
'48.5', '48.7', '48.8', '48.9', '49.2', '49.3', '49.4', '49.5',
'49.8', '49.9', '50.1', '50.2', '50.3', '50.4', '50.5', '50.6',
'50.8', '50.9', '51.0', '51.5', '51.7', '51.8', '51.9', '52.3',
'52.5', '52.7', '52.8', '52.9', '53.4', '53.5', '53.8', '53.9',
'54.0'], dtype=object)

```

Final dataset after Data preparation:

In [35]: df_brainstroke


```
Out[35]:
```

	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_
0	MALE	67.0	NO	YES	YES	PRIVATE	URBAN	
1	FEMALE	61.0	NO	NO	YES	SELF EMPLOYED	RURAL	
2	MALE	80.0	NO	YES	YES	PRIVATE	RURAL	
3	FEMALE	49.0	NO	NO	YES	PRIVATE	URBAN	
4	FEMALE	79.0	YES	NO	YES	SELF EMPLOYED	RURAL	
...
5105	FEMALE	80.0	YES	NO	YES	PRIVATE	URBAN	
5106	FEMALE	81.0	NO	NO	YES	SELF EMPLOYED	URBAN	
5107	FEMALE	35.0	NO	NO	YES	SELF EMPLOYED	RURAL	
5108	MALE	51.0	NO	NO	YES	PRIVATE	RURAL	
5109	FEMALE	44.0	NO	NO	YES	GOVERNMENT	URBAN	

5110 rows × 13 columns

```
<
In [... df_brainstroke = df_brainstroke.astype({'age': 'float', 'avg_glucose_level': 'fl
In [37]: df_brainstroke.dtypes
Out[37]:gender          object
age                    float64
hypertension          object
heart_disease         object
ever_married          object
work_type             object
Residence_type        object
avg_glucose_level     float64
bmi                   float64
smoking_status        object
brain_stroke          object
stage_of_life         object
weight_status         object
dtype: object
In [38]: df_brainstroke
```

Out[38]:

	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_
0	MALE	67.0	NO	YES	YES	PRIVATE	URBAN	
1	FEMALE	61.0	NO	NO	YES	SELF EMPLOYED	RURAL	
2	MALE	80.0	NO	YES	YES	PRIVATE	RURAL	
3	FEMALE	49.0	NO	NO	YES	PRIVATE	URBAN	
4	FEMALE	79.0	YES	NO	YES	SELF EMPLOYED	RURAL	
...
5105	FEMALE	80.0	YES	NO	YES	PRIVATE	URBAN	
5106	FEMALE	81.0	NO	NO	YES	SELF EMPLOYED	URBAN	
5107	FEMALE	35.0	NO	NO	YES	SELF EMPLOYED	RURAL	
5108	MALE	51.0	NO	NO	YES	PRIVATE	RURAL	
5109	FEMALE	44.0	NO	NO	YES	GOVERNMENT	URBAN	

5110 rows × 13 columns

<  >

```
In ... df_brain_one_hot = pd.get_dummies(data=df_brainstroke,columns=['gender','hyperten  
                                'work_type','Residence_t  
                                'weight_status'])
```

```
In [40]: df_brain_one_hot
```

Out[40]:

	age	avg_glucose_level	bmi	brain_stroke	gender_FEMALE	gender_MALE	gender_OTHER	hy
0	67.0	228.69	36.6	1	0	1	0	
1	61.0	202.21	28.1	1	1	0	0	
2	80.0	105.92	32.5	1	0	1	0	
3	49.0	171.23	34.4	1	1	0	0	
4	79.0	174.12	24.0	1	1	0	0	
...
5105	80.0	83.75	28.1	0	1	0	0	
5106	81.0	125.20	40.0	0	1	0	0	
5107	35.0	82.99	30.6	0	1	0	0	
5108	51.0	166.29	25.6	0	0	1	0	
5109	44.0	85.28	26.2	0	1	0	0	

5110 rows × 35 columns



In [41]: `df_brain_one_hot.dtypes`

```

Out[41]:age                                float64
        avg_glucose_level                  float64
        bmi                                float64
        brain_stroke                       object
        gender_FEMALE                     uint8
        gender_MALE                       uint8
        gender_OTHER                      uint8
        hypertension_NO                   uint8
        hypertension_YES                   uint8
        heart_disease_NO                  uint8
        heart_disease_YES                 uint8
        ever_married_NO                   uint8
        ever_married_YES                  uint8
        work_type_GOVERNMENT              uint8
        work_type_NEVER_WORKED            uint8
        work_type_PRIVATE                  uint8
        work_type_SELF_EMPLOYED           uint8
        work_type_STUDENT                  uint8
        Residence_type_RURAL              uint8
        Residence_type_URBAN              uint8
        smoking_status_FORMERLY SMOKED    uint8
        smoking_status_NEVER SMOKED       uint8
        smoking_status_SMOKES             uint8
        smoking_status_UNKNOWN            uint8
        stage_of_life_ADULT               uint8
        stage_of_life_CHILD               uint8
        stage_of_life_INFANT              uint8
        stage_of_life_MIDDLE AGE ADULT    uint8
        stage_of_life_SENIOR ADULT        uint8
        stage_of_life_TEEN                uint8
        stage_of_life_TODDLER             uint8
        weight_status_HEALTHY WEIGHT      uint8
        weight_status_OBESITY             uint8
        weight_status_OVERWEIGHT          uint8
        weight_status_UNDERWEIGHT         uint8
        dtype: object

```

Data Visualization

```


In [42]: '''plt.figure()
        plt.hist(df.loc[df_brainstroke['brain_stroke']=='Yes','gender'],align='mid')
        plt.xlabel('Gender')
        plt.ylabel('Count of Brain Strokes')
        plt.title('Comparing Brain Stroke in patients based on the gender category')
        plt.show()'''
        import seaborn as sns
        sns.displot(df_brainstroke, x="brain_stroke", col="gender")

```

```

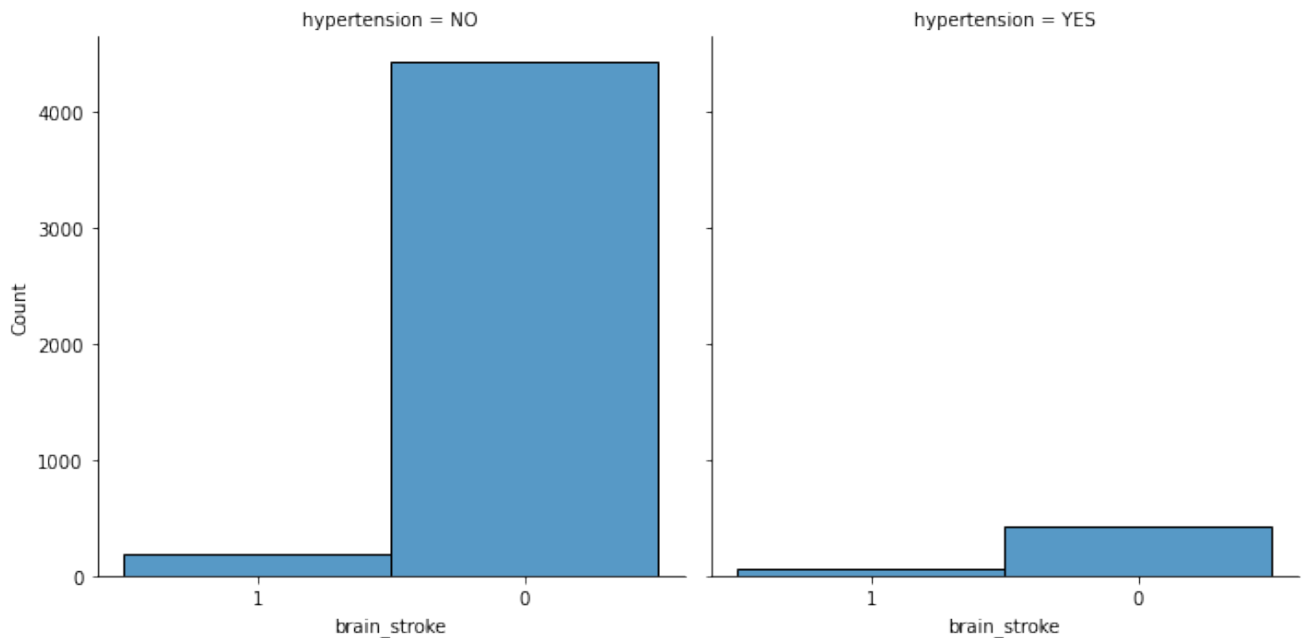
Out[42]:<seaborn.axisgrid.FacetGrid at 0x19daa4447f0>

```




```
In [4... '''plt.figure()
plt.hist(df.loc[df_brainstroke['brain_stroke']=='Yes','hypertension'],align='mi
plt.xlabel('Hypertension')
plt.ylabel('Count of Brain Strokes')
plt.title('1. Comparing Brain Stroke in patients based on hypertension')
plt.show()'''
import seaborn as sns
sns.displot(df_brainstroke, x="brain_stroke", col="hypertension")
```

Out[43]:<seaborn.axisgrid.FacetGrid at 0x19daa4442b0>



```
In [4... '''plt.figure()
plt.hist(df.loc[df_brainstroke['brain_stroke']=='Yes','heart_disease'],align='m:
plt.xlabel('Heart Disease')
plt.ylabel('Count of Brain Strokes')
plt.title('1. Comparing Brain Stroke in patients based on heart disease')
plt.show()'''
import seaborn as sns
sns.displot(df_brainstroke, x="brain_stroke", col="heart_disease")
```


Out[44]:<seaborn.axisgrid.FacetGrid at 0x19dab90d970>



As per analysis, With heart disease history have high chances of getting Brain stroke as the grph shows heart disease is directly proportional to Brain Stroke.

```
In [4... '''plt.figure()
plt.hist(df.loc[df_brainstroke['brain_stroke']=='Yes','smoking_status'],align='r
plt.xlabel('Smoking Status')
plt.ylabel('Count of Brain Strokes')
plt.title('Comparing Brain Stroke in patients based on smoking status')
plt.show()'''
import seaborn as sns
sns.displot(df_brainstroke, x="brain_stroke", col="smoking_status")
```


Out[45]:<seaborn.axisgrid.FacetGrid at 0x19daba425b0>



As per graphical representation, we can confirm that non-smokers have very little chances of getting Brain stroke

```
In [46]: '''plt.rcParams["figure.figsize"] = (20,8)
plt.hist(df.loc[df_brainstroke['brain_stroke']!='No','age'])
plt.xlabel('Age')
plt.ylabel('Count of Brain Strokes')
plt.title('Comparing Brain Stroke in patients based on age')
plt.show()'''
sns.displot(df_brainstroke, x="age", hue="brain_stroke",height=8,aspect=20/8)
```

Out[46]:<seaborn.axisgrid.FacetGrid at 0x19dab925730>



As per the graphical analysis, people over the age of 40 have high chances of getting Brain stroke

MILESTONE 3

Building a model and evaluating

```
In [47]: #importing Libraries
        # Core Libraries
        import pandas as pd
        import numpy as np
        import matplotlib.pyplot as plt

        # Sklearn processing
        from sklearn.preprocessing import MinMaxScaler
        from sklearn.model_selection import train_test_split

        # Sklearn regression algorithms
        from sklearn.linear_model import LinearRegression
        from sklearn.neighbors import KNeighborsRegressor
        from sklearn.tree import DecisionTreeRegressor

        # Sklearn regression model evaluation function
        from sklearn.metrics import mean_absolute_error

        from sklearn.neighbors import KNeighborsClassifier
        from sklearn.metrics import classification_report, confusion_matrix

        from sklearn.preprocessing import StandardScaler
        from sklearn.metrics import accuracy_score
        from sklearn.preprocessing import OneHotEncoder

        from sklearn.datasets import make_classification

In [48]: cat_columns = df_brain_one_hot.select_dtypes(include='object').keys()
        cat_columns

Out[48]: Index(['brain_stroke'], dtype='object')

In [4... # creating instance of one-hot-encoder
        enc = OneHotEncoder(handle_unknown='ignore')
```

Out[49]:

	0	1
0	0.0	1.0
1	0.0	1.0
2	0.0	1.0
3	0.0	1.0
4	0.0	1.0
...
5105	1.0	0.0
5106	1.0	0.0
5107	1.0	0.0
5108	1.0	0.0
5109	1.0	0.0

5110 rows × 2 columns

```
In [50]: # merge with main df df_brainstroke on key values
df_brainstroke = df_brainstroke.join(enc_df)
df_brainstroke
```



```
Out[50]:
```

	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_
0	MALE	67.0	NO	YES	YES	PRIVATE	URBAN	
1	FEMALE	61.0	NO	NO	YES	SELF EMPLOYED	RURAL	
2	MALE	80.0	NO	YES	YES	PRIVATE	RURAL	
3	FEMALE	49.0	NO	NO	YES	PRIVATE	URBAN	
4	FEMALE	79.0	YES	NO	YES	SELF EMPLOYED	RURAL	
...	
5105	FEMALE	80.0	YES	NO	YES	PRIVATE	URBAN	
5106	FEMALE	81.0	NO	NO	YES	SELF EMPLOYED	URBAN	
5107	FEMALE	35.0	NO	NO	YES	SELF EMPLOYED	RURAL	
5108	MALE	51.0	NO	NO	YES	PRIVATE	RURAL	
5109	FEMALE	44.0	NO	NO	YES	GOVERNMENT	URBAN	

5110 rows × 15 columns

```
< >
In... df_brainstroke['gender'] = df_brainstroke['gender'].map(dict(zip(['MALE', 'FEMALE'],
df_brainstroke['hypertension'] = df_brainstroke['hypertension'].map(dict(zip(['YES', 'NO'],
df_brainstroke['heart_disease'] = df_brainstroke['heart_disease'].map(dict(zip(['NO', 'YES'],
df_brainstroke['ever_married'] = df_brainstroke['ever_married'].map(dict(zip(['YES', 'NO'],
df_brainstroke['work_type'] = df_brainstroke['work_type'].map(dict(zip(['PRIVATE', 'GOVERNMENT',
'NEVER_WORKED'], [0,1,2,3,4])))
df_brainstroke['Residence_type'] = df_brainstroke['Residence_type'].map(dict(zip(['URBAN', 'RURAL'],
df_brainstroke['smoking_status'] = df_brainstroke['smoking_status'].map(dict(zip(['SMOKER', 'NON-SMOKER'],
df_brainstroke['stage_of_life'] = df_brainstroke['stage_of_life'].map(dict(zip(['TODDLER', 'CHILD'], [0,1,2,3,4,5,6])))
df_brainstroke['weight_status'] = df_brainstroke['weight_status'].map(dict(zip(['LOW', 'NORMAL', 'HIGH'],
In [52]: #df_brain = df_brain.dropna()
y = df_brain_one_hot["brain_stroke"]
X = df_brain_one_hot.drop('brain_stroke', axis=1)
In [53]: # Split dataset into random train and test subsets:
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20)
In [54]: #pip install imbalanced-Learn
```

```
In [56]: # summarize the new class distribution
        from collections import Counter
        counter = Counter(y_train)
        print(counter)
        counter = Counter(y_test)
        print(counter)
        counter = Counter(y_res)
        print(counter)
```

```
Counter({'0': 3895, '1': 193})
Counter({'0': 966, '1': 56})
Counter({'0': 3895, '1': 3895})
```

```
In [57]: # Standardize features by removing mean and scaling to unit variance:
        scaler = StandardScaler()
        scaler.fit(X_train)
```

```
Out[57]:StandardScaler()
```

```
In [58]: X_train = scaler.fit_transform(X_train)
        X_test = scaler.fit_transform(X_test)
        X_res = scaler.fit_transform(X_res)
```

2. KNN classifier:

```
In [59]: # Use the KNN classifier to fit data:
        classifier = KNeighborsClassifier()
        classifier.fit(X_res, y_res)
```

```
Out[59]:KNeighborsClassifier()
```

```
In [60]: # Predict y data with classifier:
        y_predict = classifier.predict(X_test)
```

```
In [61]: # Print results:
        print(confusion_matrix(y_test, y_predict))
        print(classification_report(y_test, y_predict))
```

```
[[937  29]
 [ 49   7]]
```

	precision	recall	f1-score	support
0	0.95	0.97	0.96	966
1	0.19	0.12	0.15	56
accuracy			0.92	1022
macro avg	0.57	0.55	0.56	1022
weighted avg	0.91	0.92	0.92	1022

```
In [62]: knn_score = accuracy_score(y_test, y_predict, normalize=False)
        knn_score
```

```
Out[62]:944
```

```
In [63]: # print the shapes of the new X objects
        print(X_train.shape)
        print(X_test.shape)
```

```

(4088, 34)
(1022, 34)
In [64]: # print the shapes of the new y objects
        print(y_train.shape)
        print(y_test.shape)

(4088,)
(1022,)
In [65]: #Importing Libraries
        from sklearn.linear_model import LogisticRegression

In [66]: #train the model on the training set
        logreg = LogisticRegression()
        logreg.fit(X_train, y_train)

Out[66]:LogisticRegression()
In [67]: import sklearn.metrics as metrics
        #make predictions on the testing set
        y_pred = logreg.predict(X_test)

        print(accuracy_score(y_test, y_pred))

0.9452054794520548
In [68]: #Repeat for KNN with K=5:
        knn = KNeighborsClassifier(n_neighbors=5)
        knn.fit(X_train, y_train)
        y_pred = knn.predict(X_test)
        print(metrics.accuracy_score(y_test, y_pred))

0.9442270058708415
In [69]: #Repeat for KNN with K=1:
        knn = KNeighborsClassifier(n_neighbors=5)
        knn.fit(X_train, y_train)
        y_pred = knn.predict(X_test)
        print(metrics.accuracy_score(y_test, y_pred))

0.9442270058708415
In [70]: # try K=1 through K=35 and record testing accuracy
        k_range = range(1, 36)

        # We can create Python dictionary using [] or dict()
        scores = []

        # We use a loop through the range 1 to 36
        # We append the scores in the dictionary
        for k in k_range:
            knn = KNeighborsClassifier(n_neighbors=k)
            knn.fit(X_train, y_train)
            y_pred = knn.predict(X_test)
            scores.append(metrics.accuracy_score(y_test, y_pred))

        print("Accuracy Score of KNN Classifier: " + str(scores))

```


Accuracy Score of KNN Classifier: [0.9187866927592955, 0.9422700587084148, 0.9344422700587084, 0.9461839530332681, 0.9442270058708415, 0.9452054794520548, 0.9442270058708415, 0.9452054794520548, 0.9452054794520548, 0.9442270058708415, 0.9452054794520548]

```
In [71]: # import Matplotlib (scientific plotting library)
import matplotlib.pyplot as plt

# allow plots to appear within the notebook
%matplotlib inline

# plot the relationship between K and testing accuracy
# plt.plot(x_axis, y_axis)
plt.plot(k_range, scores)
plt.xlabel('Value of K for KNN')
plt.ylabel('Testing Accuracy')
```

Out[71]:Text(0, 0.5, 'Testing Accuracy')



Summary:

As per the classification regression and KNeighborsRegressoion after evaluating the accuracy of the training model we got the output around 0.95

Decision Tree Classifier:

```
In [72]: #Import Libraries
from sklearn.tree import DecisionTreeClassifier

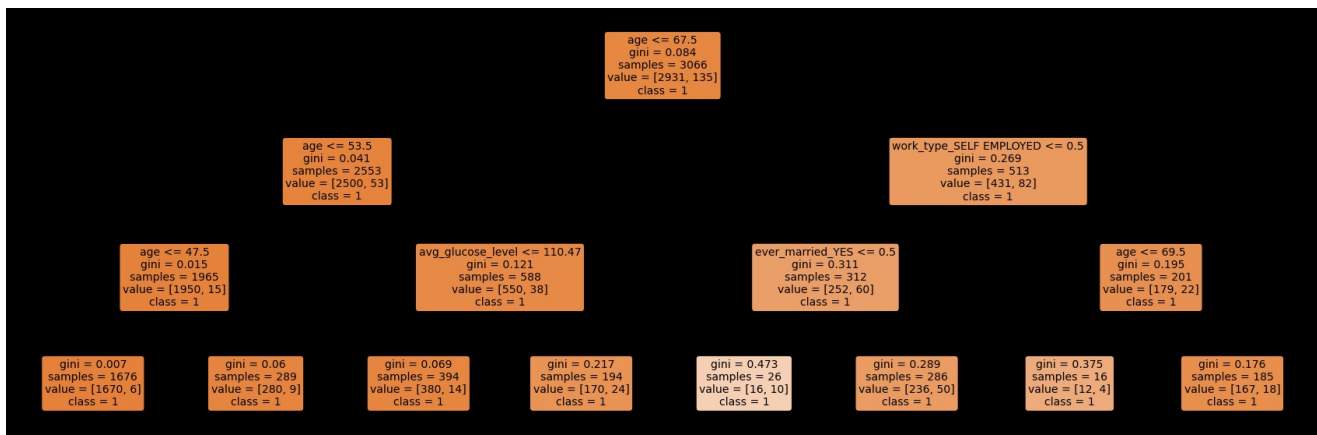
In [73]: #save the feature name and target variables
feature_names = X.columns
labels = y.unique()
#split the dataset
from sklearn.model_selection import train_test_split
X_train, test_x, y_train, test_lab = train_test_split(X,y,
```

```
In [74]: clf = DecisionTreeClassifier(max_depth =3, random_state = 42)
        clf.fit(X_train, y_train)
```

```
Out[74]:DecisionTreeClassifier(max_depth=3, random_state=42)
```

```
In [75]: #import relevant packages
        from sklearn import tree
        import matplotlib.pyplot as plt
        #plt the figure, setting a black background
        plt.figure(figsize=(30,10), facecolor = 'k')
        #create the tree plot
        a = tree.plot_tree(clf,
                           #use the feature names stored
                           feature_names = feature_names,
                           #use the class names stored
                           class_names = labels,
                           rounded = True,
                           filled = True,
                           fontsize=14)


        #show the plot
        plt.show()
```



```
In [76]: test_pred_decision_tree = clf.predict(test_x)
```

```
In [77]: #import the relevant packages
        from sklearn import metrics
        import seaborn as sns
        import matplotlib.pyplot as plt
        #get the confusion matrix
        confusion_matrix = metrics.confusion_matrix(test_lab,
                                                    test_pred_decision_tree)

        #turn this into a dataframe
        matrix_df = pd.DataFrame(confusion_matrix)
        #plot the result
        ax = plt.axes()
        sns.set(font_scale=1.3)
        plt.figure(figsize=(10,7))
        sns.heatmap(matrix_df, annot=True, fmt="g", ax=ax, cmap="magma")
        #set axis titles
        ax.set_title('Confusion Matrix - Decision Tree')
        ax.set_xlabel("Predicted label", fontsize =15)
```



<Figure size 720x504 with 0 Axes>

```
In [... print("Accuracy Score of Decision Tree Classifier: " + str(accuracy_score(test_]
```

Accuracy Score of Decision Tree Classifier: 0.9442270058708415

Random Forest Classifier

```
In [... #Split the data
```

```
    training, testing, training_labels, testing_labels = train_test_split(X, y, test_
```

```
In [8... # Normalize the data
```

```
    sc = StandardScaler()
```

```
    normed_train_data = pd.DataFrame(sc.fit_transform(training), columns = X.column
```

```
    normed_test_data = pd.DataFrame(sc.fit_transform(testing), columns = X.columns)
```

```
In [81]: #Building Random Forest model
```

```
    from sklearn.ensemble import RandomForestClassifier
```

```
    from sklearn.metrics import confusion_matrix
```

```
    clf=RandomForestClassifier()
```

```
    clf.fit(training, training_labels)
```

```
Out[81]:RandomForestClassifier()
```

```
In [82]: preds = clf.predict(testing)
```

```
In [83]: print(clf.score(testing, testing_labels))
```

0.9374021909233177

```
In [... print("Accuracy Score of Random Forest Classifier: " + str(accuracy_score(testi
```

Accuracy Score of Random Forest Classifier: 0.9374021909233177

```
In [85]: confusion_matrix(testing_labels, preds)
```

```
Out[85]:array([[1197,    1],
               [   79,    1]], dtype=int64)
```

In []: